



# Tactics of Adversarial Attack on Deep Reinforcement Learning Agents

Yen-Chen Lin, Zhang-Wei Hong, Yuan-Hong Liao, Meng-Li Shih, Ming-Yu Liu, Min Sun

- Prakash K Naikade, [prna00001@stud.uni-saarland.de](mailto:prna00001@stud.uni-saarland.de)

# Paper's Contribution

- Empirical evidences shows studying the adversarial attack on deep RL agents is critical
- Strategically-Timed Attack
  - Attacking at critical moments in an episode to minimize the agent's reward
- Enchanting Attack
  - Luring the agent to a designated target/dangerous state

- Reinforcement Learning = RL
- Deep Neural Network = DNN

# Why Adversarial Reinforcement Learning?

- DNNs are ideal function approximators for classical RL algorithms
- DNNs are vulnerable to the adversarial example attack
- Adversarial attack on deep RL agents is different from adversarial attack on classification system
  - Reducing rewards or luring agent to dangerous state Vs Reducing classification accuracy
- Real world adversarial examples and examples crafted by adversary
- Critical to understand these vulnerabilities to failproof the RL algorithms used in mission-critical tasks

# Why New Tactics?

- Existing tactics
  - *Uniform Attack* [Huang et al., 2017]
  - Ignores the fact that observations are correlated
  - Attack at every time step
  - More prone to detection
- Goal should be
  - Attack at selective time steps
  - At a time step where it could be more effective
  - E.g. Ping Pong
    - When the ball is away from paddle - no need to attack
    - When the ball is close to the paddle - attacking could cause the dropping of the ball

# Why new tactics?

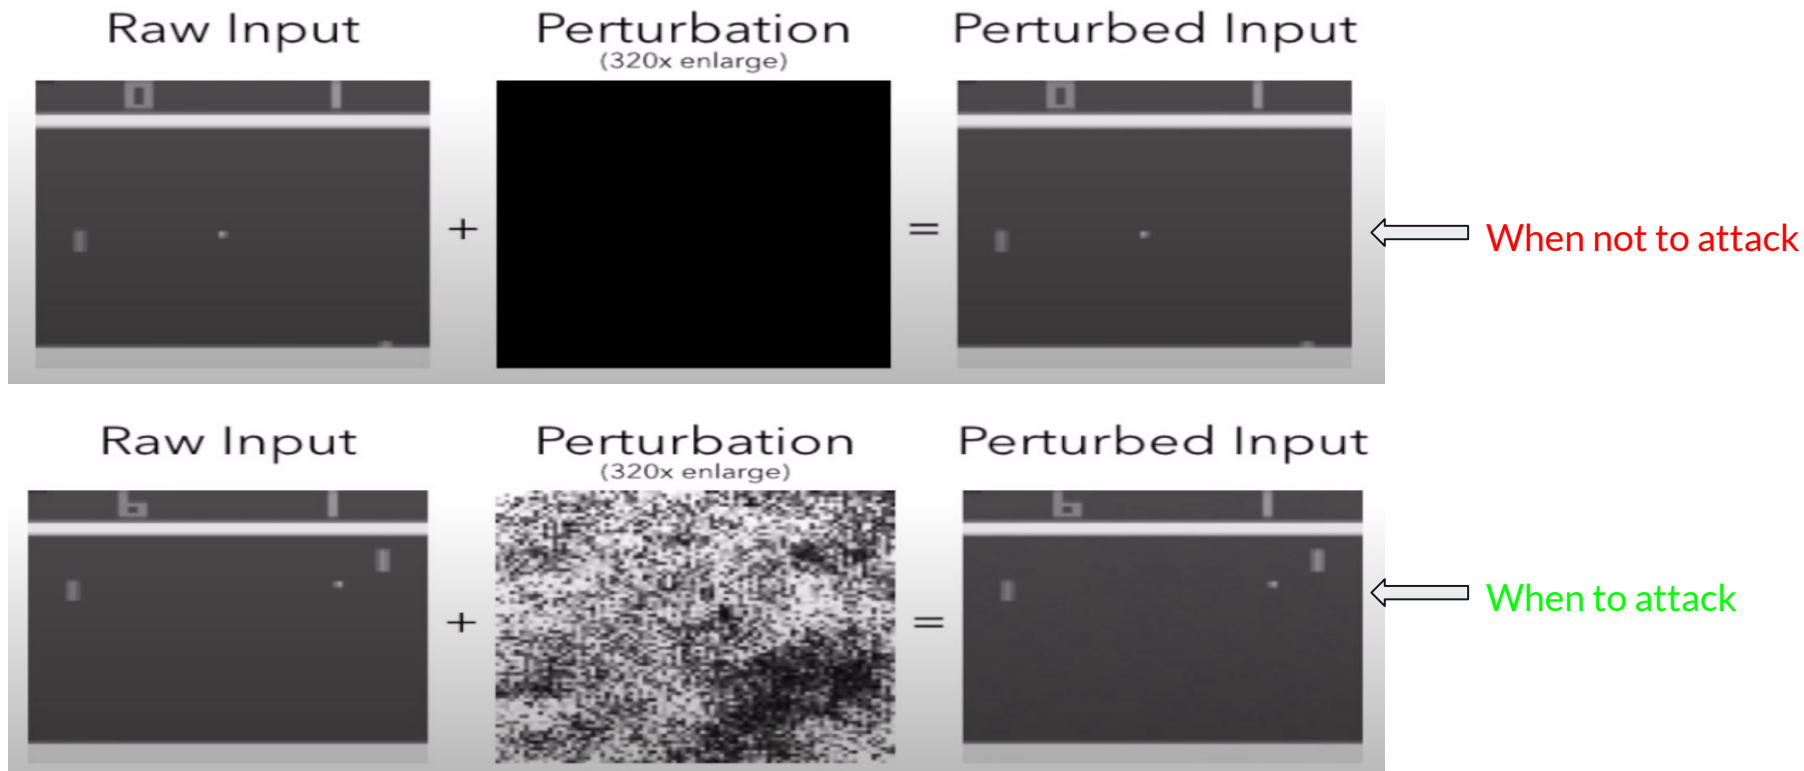


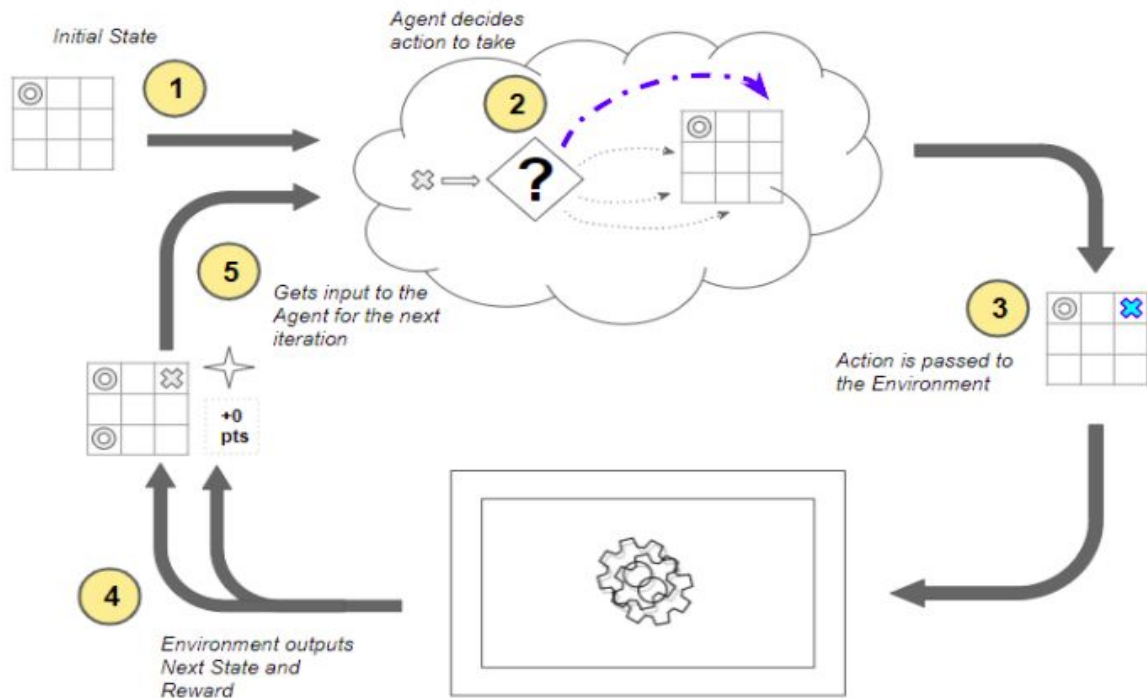
Fig 1 - When to attack

# Two Novel Attacks

- Strategically-Timed Attack
  - Aims to reduce the reward with as fewer adversarial examples as possible
  - Adversarial example is only used when the attack is expected to be effective
  - Achieves same effect as the uniform attack by attacking four times less
- Enchanting Attack
  - Maliciously misguiding agent to a target state
  - Success rate > 70% in attacking agents

# Reinforcement Learning

Let's structure RL problem as a Markov Decision Process (MDP),



MDP has 5 components,

- Agent,
- Environment,
- State,
- Action,
- Reward

Fig 2 - How the MDP works (Image by Ketan Doshi)

# Reinforcement Learning

Let's structure RL problem as a Markov Decision Process (MDP),

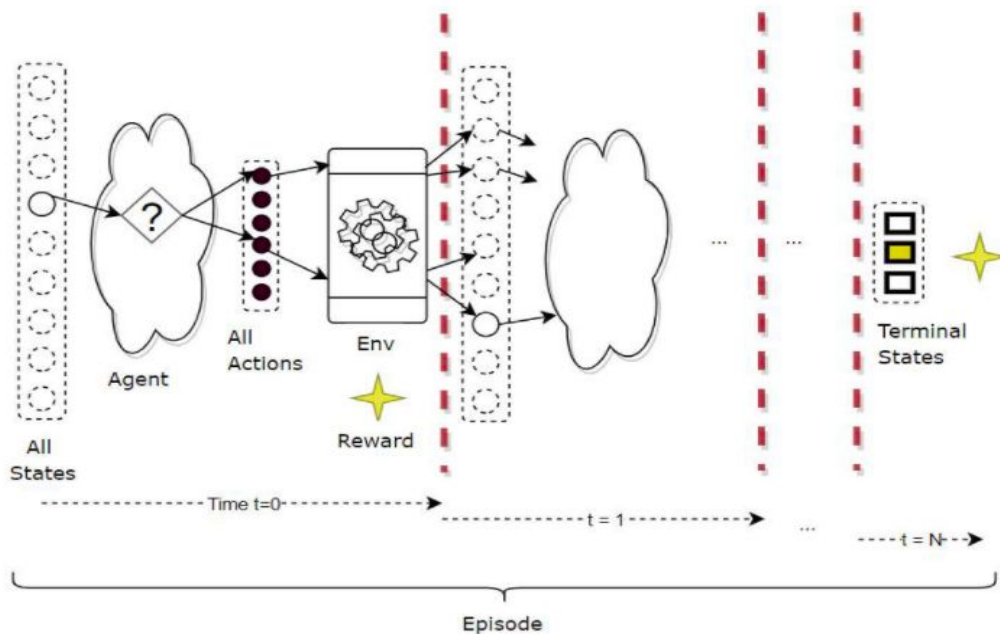


Fig 3 - An MDP iterates over a sequence of time steps  
(Image by Ketan Doshi)

- At each time step,
  - Agent performs an action based on the observation of the environment
  - For maximizing the accumulated future rewards
  - Action determination is through a policy  $\pi$
- Goal of RL algorithm at time step  $t$ ,
  - learn a policy that maximizes the accumulated future rewards,  
$$\mathcal{R}_t = r_t + r_{t+1} + \dots + r_{\mathcal{L}}$$
  - policy  $\pi$  is learned through DNN
- Attack by,
  - Perturbing the observations by using crafted adversarial example
  - Sequence of adversarial example to lure agent to dangerous state



# Strategically-Timed Attack

- Using our observation, we could minimize the expected return at first time step, by following optimization problem,

$$\begin{aligned} \min_{b_1, b_2, \dots, b_L, \delta_1, \delta_2, \dots, \delta_L} \quad & R_1(\bar{s}_1, \dots, \bar{s}_L) \\ \bar{s}_t = s_t + b_t \delta_t & \quad \text{for all } t = 1, \dots, L \\ b_t \in \{0, 1\}, & \quad \text{for all } t = 1, \dots, L \\ \sum_t b_t \leq \Gamma & \end{aligned} \tag{1}$$

Where,

$\{s_1, \dots, s_L\}$  are sequence of observations or states in an episode,  
 $\{\delta_1, \dots, \delta_L\}$  are sequence of perturbations,  
 $\mathcal{R}_1$  is the expected return at the first time step,  
 $b_1, \dots, b_L$  denote when to attack.

- Mixed integer programming problem
- Problem size grows exponentially with L
- Divide the problem (1) into “When-to-Attack” ( $b_1, \dots, b_L$ ) and “How-to-Attack” ( $\{\delta_1, \dots, \delta_L\}$ )

- Relative action preference function  $c$  to compute  $b_1, \dots, b_L$

- For *policy gradient-based methods* (e.g. A3C algorithm),

$$c(s_t) = \max_{a_t} \pi(s_t, a_t) - \min_{a_t} \pi(s_t, a_t)$$

Where,  $\pi$  is policy network which maps a state-action pair  $(s_t, a_t)$  to a probability

- For *value-based methods* (e.g. DQN algorithm),

$$c(s_t) = \max_{a_t} \frac{e^{\frac{Q(s_t, a_t)}{T}}}{\sum_{a_k} e^{\frac{Q(s_t, a_k)}{T}}} - \min_{a_t} \frac{e^{\frac{Q(s_t, a_t)}{T}}}{\sum_{a_k} e^{\frac{Q(s_t, a_k)}{T}}}$$

- Represents criticality to perform specific action to increase return
- Performs attack :

$$b_t = 1 \text{ when } c(s_t) \geq \beta$$

$\beta$  controls the number of attacks to be performed

- Change most preferred action to least preferred to reduce accumulated reward
- Cue: Output of a trained deep RL agent to craft effective adversarial example by adding perturbation
- Adversarial example crafting method [Carlini and Wagner, 2016]
  - Approximating following optimization problem,

$$\min_{\delta} \mathcal{D}_I(x, x + \delta), \text{ subject to } f(x) \neq f(x + \delta)$$

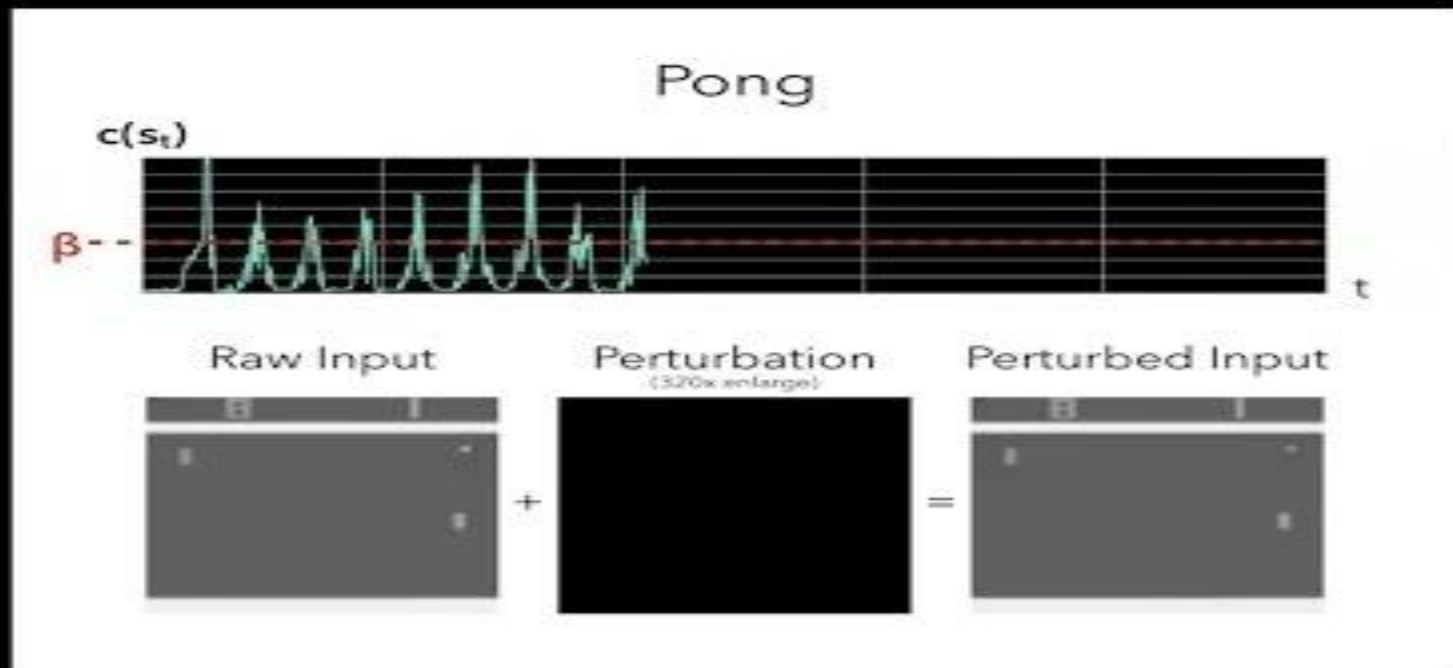
where,  $x$  is an image,

$f$  is DNN,

$\mathcal{D}_I$  is an image similarity metric which finds minimal perturbation  $\delta$

# Strategically-Timed Attack

- Example: ([http://yenchelin.me/adversarial\\_attack\\_RL/](http://yenchelin.me/adversarial_attack_RL/))



# Enchanting Attack

- To lure agent from current state  $s_t$ , at time step  $t$ , to a specified target state  $s_g$ , after  $H$  steps
- Need to craft a series of adversarial example  $s_{t+1} + \delta_{t+1}, \dots, s_{t+H} + \delta_{t+H}$
- Problem divided into two subtasks
- First Subtask:
  - Assume that adversary has full control of the agent to take arbitrary actions at each step
  - Task reduced to planning a sequence of actions for reaching the target state
  - By using online planning algorithm
- Second Subtask:
  - Craft example  $s_t + \delta_t$
  - To force an agent to take the first action of planned action sequence
  - Gradually create adversarial examples in order to persuade an agent to take planned action sequence  $s_{t+1} + \delta_{t+1}, \dots, s_{t+H} + \delta_{t+H}$  [Carlini and Wagner, 2016]

- “Future state prediction” + “Sampling-based action planning”
- *Future state prediction and evaluation:*
  - Next video-frame prediction model  $M$  to predict a future state given a sequence of actions ,

$$s_{t+H}^M = M(s_t, A_{t:t+H}) \quad [\text{Oh et al., 2015}]$$

Where,

$A_{t:t+H} = \{a_t, \dots, a_{t+H}\}$  is the given sequence of  $H$  future actions ,  
beginning at step  $t$ ,  
 $s_t$  is the current state,  
 $s_{t+H}^M$  is the predicted future state

- Success of the attack based on distance between  $s_g$  and  $s_{t+H}^M$  ,  
which is given by  $D(s_g, s_{t+H}^M)$  ..... ( $L_2$ - norm)

- Future state prediction + Sampling-based action planning
- *Sampling-based action planning:*
  - Sampling-based cross-entropy method to compute a sequence of actions [Rubinstein and Kroese, 2013]
  - Sample  $N$  action sequences of length  $H : \{A_{t:t+H}^n\}_{n=1}^N$
  - Rank each of them based on the distance between the final state obtained after performing the action sequence and the target state  $s_g$
  - Keep the best  $K$  action sequences and refit our categorical distributions to them

# Flow of Enchanting Attack

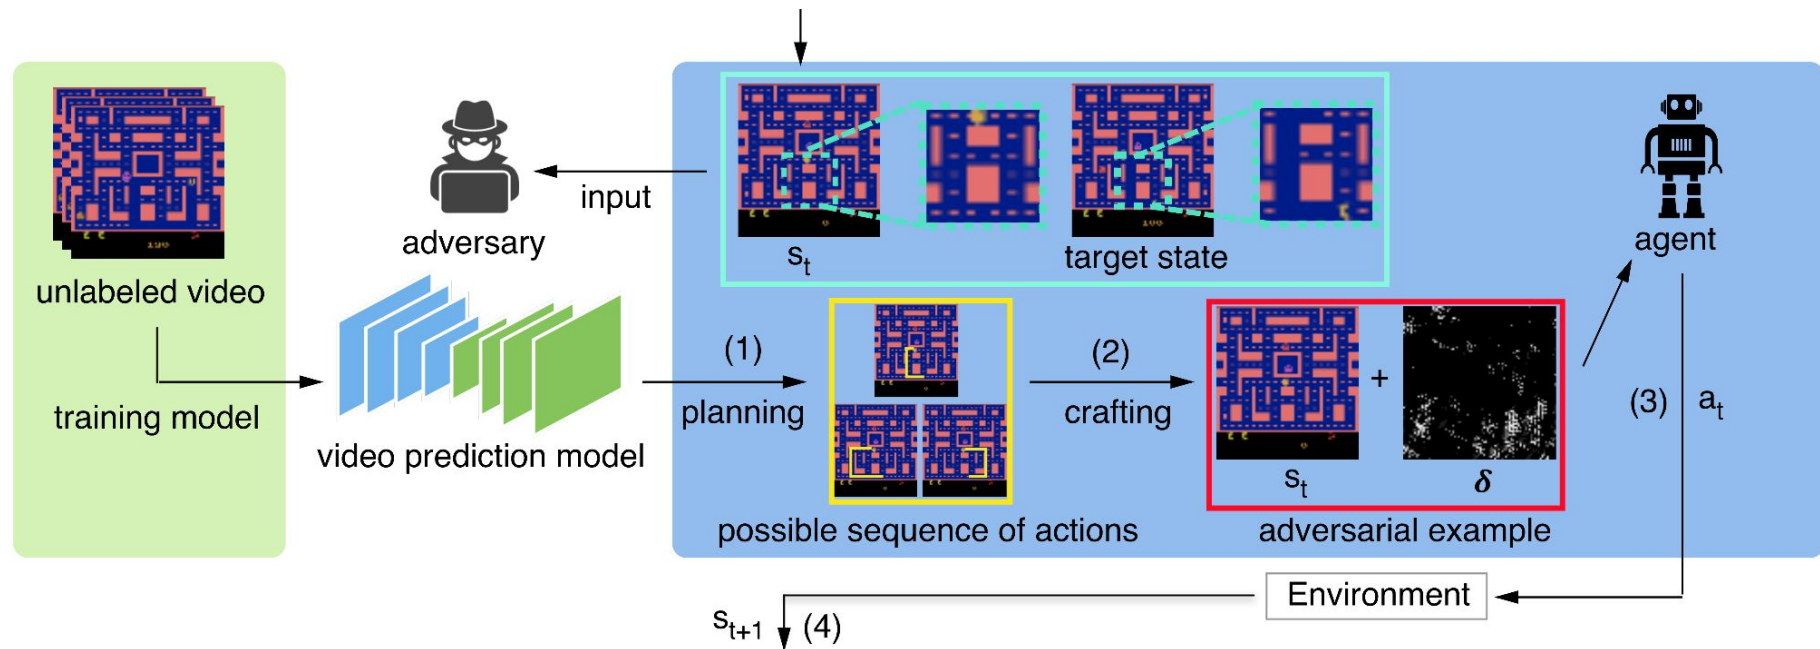


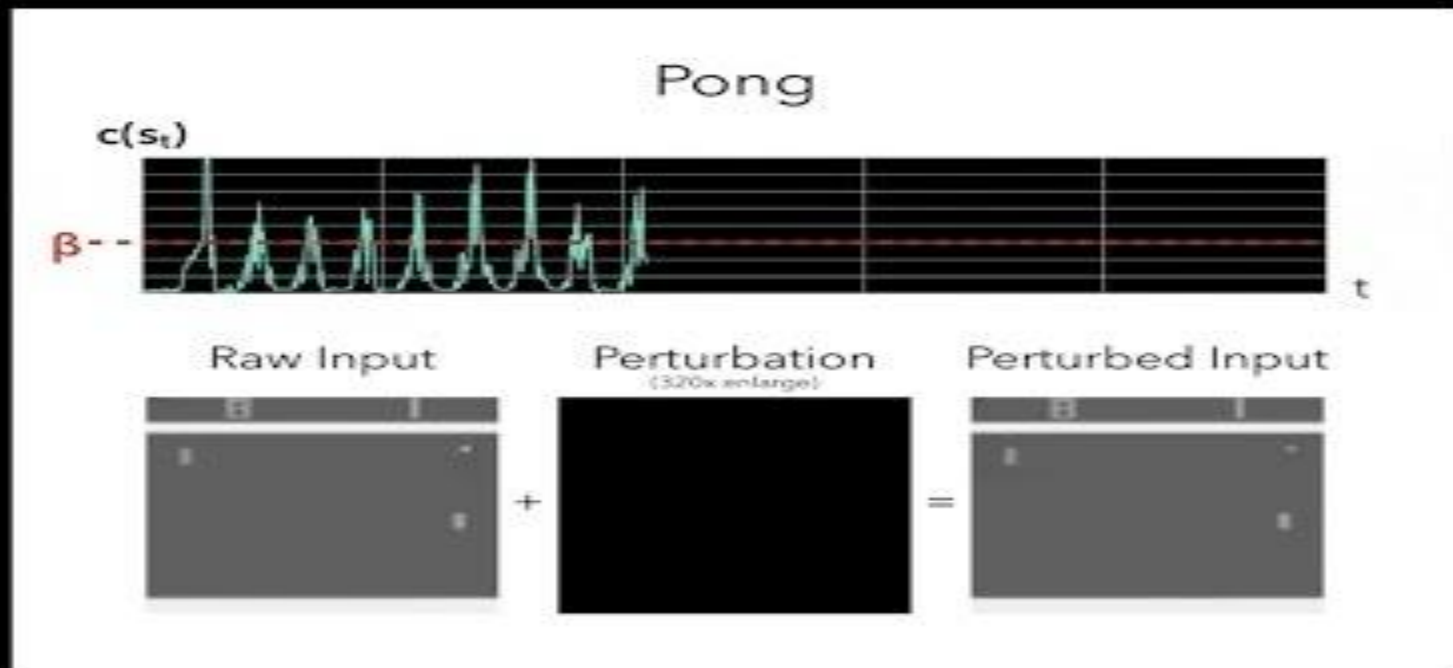
Fig 4 - Illustration of Enchanting Attack on Ms.Pacman.

- Instead of directly crafting the next adversarial example with target action  $a_{t+1}$ , start another enchanting attack at state  $s_{t+1}$  for robustness



# Enchanting Attack

- Example: ([http://yenchenlin.me/adversarial\\_attack\\_RL/](http://yenchenlin.me/adversarial_attack_RL/))



# Experimental Results

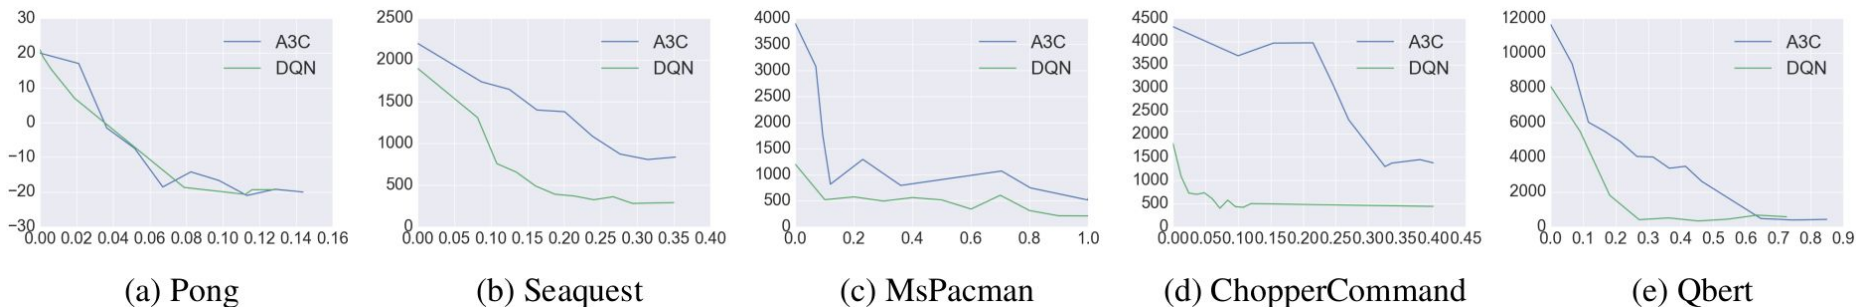


Fig 5 - Accumulated reward (y-axis) v.s. Portions of time steps the agent is attacked (x-axis) of Strategically-timed Attack

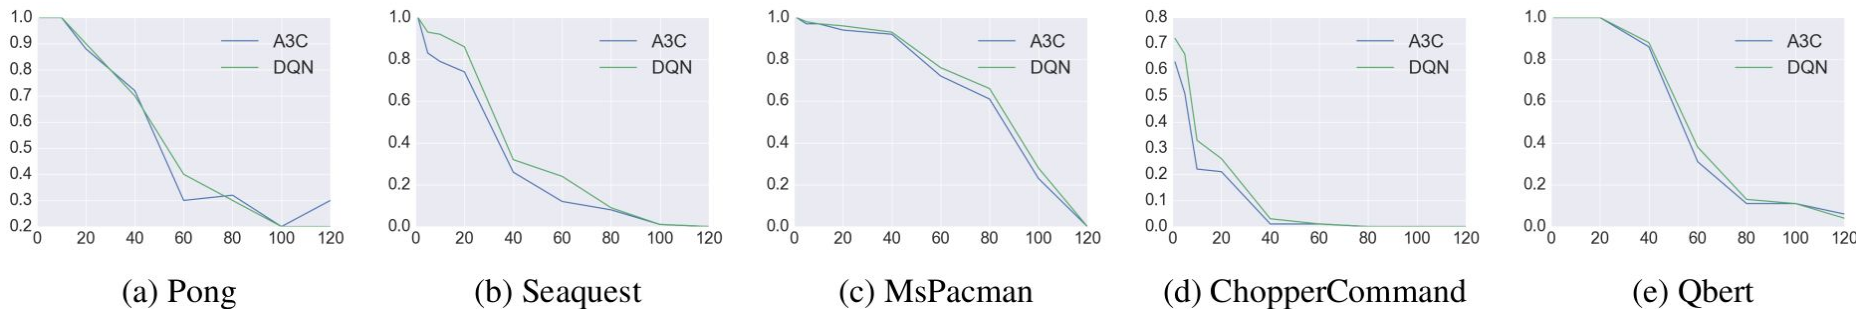


Fig 6 - Success rate (y-axis) v.s.H steps in the future (x-axis) for Enchanting Attack

# Experimental Results

- Strategically-timed attack can achieve same effect as the uniform attack by attacking just 25% of the time steps in an episode
- Enchanting attack forces agent toward maliciously defined target state in 40 steps with more than 70% success rate in 3 out of 5 games
- Enchanting attack was less effective on Seaquest and ChopperCommand
- Both games include multiple random enemies and therefore trained video prediction models were less accurate
- An agent trained using the DQN algorithm was more vulnerable than an agent trained with the A3C algorithm in most games except Pong
- Stronger deep RL agent may be more robust to the adversarial attack

# Strengths

- In comparison to Huang et al., 2017, this is an innovative use of a "classic test time" adversarial attacks
- The overall experimental design is sound, and it covers all potential scenarios
- The research topic is interesting, and the contributions are substantial

# Weaknesses

- The strategically-timed attack technique only considered the attack effect on one time step, ignoring the impact on the following states and actions, i.e., not considering the final end goal
- In an enchanting attack, accurate prediction and enforcement of future states and actions is difficult, and thus this approach suffers from a low attack success rate
- The experimental evaluation in general and in-terms of performance comparison with Huang et al., 2017 is weak and vague
- The paper's flaw is that the problem isn't well-motivated, and the authors jumped right into the content without explaining the publications they utilized to develop these adversarial tactics

# Possible Improvements



- This paper requires a more thorough overview of the relevant work on which it is based
- More fluid narrative and structure

# Possible Extensions

- Developing defenses against these kind of adversarial attacks
- Develop a more sophisticated strategically-timed attack strategy
- enhancing the generative model's video prediction accuracy in order to increase the success rate of enchantment attacks in more complex environments



Thank You !

Questions Please !